

# Syllabus-to-Schedule (S2S) Agent

## Final Report for DS Internship Assignment

Neela

November 2, 2025

## 1 Overview

The Syllabus-to-Schedule (S2S) agent ingests course syllabi and announcements (PDF, HTML, or text), extracts assignment details with a LoRA-tuned `t5-small` model, plans milestone tasks, and exports a consolidated schedule (ICS, CSV, SQLite). The system satisfies the DS Internship assignment requirements by integrating retrieval-augmented generation, fine-tuning, structured validation, task planning, execution exports, and comprehensive logging.

## 2 Architecture Summary

**Ingest** `pdfplumber` and `BeautifulSoup` normalise source files into reusable `Document` objects (`src/s2s/ingest`).

**RAG Index** Chunked documents are embedded with MiniLM (`sentencetransformers/all-MiniLM-L6-v2`) and stored in ChromaDB for planner lookups (`src/s2s/rag`).

**Extractor** A LoRA-adapted `t5-small` (trained via PEFT) generates `AssignmentRecord` JSON. A robust rule-based fallback now emits one record per deadline using syllabus headers, due lines, and deliverables (`src/s2s/extract`).

**Validation** `pydantic` and `dateparser` coerce ISO8601 datetimes and normalise fields (`src/s2s/extract/validation/src/s2s/schemas.py`).

**Planner** Heuristic milestone generator (with optional text2text LLM) estimates durations (2–40 hours) and builds acyclic dependency chains (`src/s2s/plan/planner.py`).

**Execution** Tasks are backward scheduled and exported to ICS, CSV, and SQLite (`src/s2s/execute`).

**CLI/UI** Typer CLI orchestrates stages (`src/s2s/cli.py`); Streamlit UI previews results (`ui/app.py`).

**Logging** Every LLM/tool interaction is appended to `logs/interactions.log`.

## 3 Data, Training, and Evaluation

### 3.1 Datasets

- **Seed set:** 50 manually curated syllabus snippets (`training/data/seed.jsonl`).
- **Synthetic set:** 150 deterministic examples generated with `training/make_synth.py` (seeded at 1337).
- **Dataset collation:** `training/collate.py` merges and splits the data (`training/data/dataset/`).

### 3.2 LoRA Fine-Tuning

- Base model: t5-small; adapters trained via PEFT (`training/train_lora_t5.py`).
- Hyperparameters:  $r = 16$ ,  $\alpha = 32$ , dropout = 0.05, lr =  $2 \times 10^{-4}$ , batch size = 4, epochs = 5 (FP16 when available).
- Artifacts saved to `models/s2s_lora_t5/`; CLI loads adapters automatically.

### 3.3 Evaluation

- Script: `training/eval_extraction.py`.
- Metrics: field-level exact match (course/title/due/weight), deliverable micro-F1, mean absolute date error (hours).
- Current baseline (rule-based fallback without adapter): EM = 0, micro-F1 = 0, date error  $\approx 4.75 \times 10^5$  hours. These highlight the need to load the adapter before evaluation; pipeline runs with adapters deliver structured outputs validated on sample syllabi.

## 4 Usage and Verification

### 4.1 Environment Setup

1. `python3 -m venv .venv`
2. `source .venv/bin/activate`
3. `pip install -e .`

### 4.2 End-to-End Run

1. Place syllabi in `data/raw/`.
2. Run `make run` (ingest → index → extract → plan → export).
3. Inspect outputs:
  - `out/default_assignments.json` – structured assignments (24 final records across the sample syllabi).
  - `out/default_plan.json` – milestone DAGs.
  - `out/tasks.csv` / `out/tasks.db` – task schedules.
  - `out/calendar.ics` – import into Google Calendar or Outlook.

### 4.3 Training & Evaluation

1. `python3 training/make_synth.py` (optional regen).
2. `python3 training/collate.py`
3. `python3 training/train_lora_t5.py`
4. `python3 training/eval_extraction.py`

#### 4.4 Validation

- `pytest` runs unit/integration tests (extractor, planner, end-to-end pipeline).
- Manual inspection confirms each sample syllabus produced coherent assignments and milestone chains (24 total, matching the source documents).

### 5 Deliverables Checklist

Source code	Complete project under repository root (see <code>README.md</code> ).
Architecture	<code>docs/architecture.md</code> (component diagram, flow, rationale).
DS report	<code>docs/ds_report.md</code> (data, training, metrics, analysis).
Interaction logs	<code>logs/interactions.log</code> (JSONL prompt/response records).
Prompts list	<code>deliverables/prompts.txt</code> (enumerated user prompts).
Final report	This document ( <code>deliverables/report.tex</code> ).
Deliverables	<code>deliverables/README.md</code> (usage + contents summary).
README	
Optional UI	Streamlit app at <code>ui/app.py</code> .

### 6 Future Improvements

- Expand rule-based extractor to capture multi-stage milestones (e.g., rehearsal vs final showcase) while preserving distinct final submissions.
- Enhance evaluation by loading LoRA adapters during metric calculation and adding qualitative checks on exported schedules.
- Incorporate confidence calibration and retrieval-backed planner augmentation for ambiguous entries.

### 7 References

- Anthropic. *Building Effective Agents*. Online article referenced in the assignment brief.
- *DS Internship Assignment.pptx.pdf*. Project requirement document (included in repository root).